


One document discipline governs agent work from intent to release

An evaluation briefing for engineering leaders whose teams build
with AI agents



2 Agents now write most production code without a shared practice

- Most professional developers ship agent-written code every day
- The output is locally good and globally inconsistent
- The document disciplines human teams built never reached the agents
- Every practice era gets its equivalent of test-driven development



how agent-written code became the default before any discipline caught up

3 Without that practice, agent code drifts from its governing documents

- Specs, designs, and tests fall out of sync because nobody walks them between sessions
- A feature ships whose design contradicts a decision recorded months earlier
- Each new conversation re-improvises the workflow, so reviews that should happen do not

1

Drift

specs and code
diverge between
sessions

2

Contradiction

a design ignores an
earlier decision

3

Lost context


every session starts
from zero

4 HELIX answers with 53 document templates across a seven-activity loop


Activity	What the templates produce
Discover	vision, business case, market and competitive analysis
Frame	requirements, features, user stories, concerns, principles
Design	architecture, decisions, contracts, solution and technical designs
Test	test plans and quality expectations, written before code
Build	implementation plans against those tests
Deploy	checklists, runbooks, monitoring, release notes
Iterate	metrics, backlog, iteration plans, status reports, deliverables

5 One skill routes twenty-one modes over those documents

- A single skill installs on the agent runtime and reads the same documents
- It routes each request to one mode: frame, align, evolve, review, present, and more
- Alignment is the required mode; it finds drift and hands back a plan
- The skill body is portable: no runtime commands, so it loads on any host




6 The alignment check reads those documents before anyone codes



the step nobody does by hand today

7 Fifty concerns carry shared practices into every document

- A concern is a cross-cutting rule set: accessibility, security, verification, data quality
- Each concern names the practices every downstream document and task inherits
- Principles settle tensions the same way each time, starting with spec is the contract
- Exclusive slots hold one choice each: one frontend framework, one runtime, one auth provider



8 Every shared practice ends in evidence a validator can check

The claim

A test covers a requirement

The work is complete

Someone promises a metric


The evidence HELIX requires

- ✓ The test cites the requirement it exercises, by id
- ✓ A validator ran and the record names the command and the guard branches
- ✓ The metric definition and its dashboard row exist

Verdict: a claim without evidence blocks the work as a phantom claim

9 HELIX checks your documents and leaves your runtime and stack alone

- No command-line tool, tracker, queue, or execution engine ships with it
- It installs on six hosts today, from Claude Code to Databricks Genie
- Your technology choices stay yours; HELIX imposes none
- How much an agent does without asking is a three-position dial you set



10

Those documents also become the decks and status reports people read

- Roadmaps, iteration plans, and status reports come from the same governed documents
- A status report cites evidence per claim; a client cut never contradicts it
- A deck or brief projects the documents in a human voice, every number sourced
- This deck came out of that pipeline, with its own gates and inspection

1

Status reports

evidence per claim
against the plan

2

Decks and briefs

the human-facing
projection

3

Release notes

from the deploy
documents

11 Pilot HELIX for one quarter and count what the documents catch

Pilot HELIX for one quarter and count what the documents catch

	Step	Owner	Date
1	Pick one active project with an agent already in the loop, by the end of this week	_____	_____
2	Adopt the templates and run one alignment check on its documents in the first sprint	_____	_____
3	Review the findings together after four weeks against the bar of under 3 drift findings per check	_____	_____

Without a trial the practice stays improvised and drift keeps surfacing in review

Sources

Id	Claim or figure	Governing artifact and section
S1	A methodology and artifact catalog for AI-assisted teams: templates plus one skill	docs/helix/00-discover/product-vision.md#mission-statement
S2	AI-assisted development went from novelty to default; most developers ship AI-touched code daily	docs/helix/00-discover/product-vision.md#why-now
S3	Three failure modes: artifact drift, inconsistent agent behavior, lossy session boundaries	docs/helix/01-frame/prd.md#problem
S4	Templates for 53 artifact types plus a single routing skill; no CLI, tracker, or execution loop	docs/helix/01-frame/prd.md#summary
S5	The seven activities, the artifact loop, and the authority hierarchy	workflows/README.md#activities , workflows/README.md#authority-hierarchy , workflows/artifact-hierarchy.md
S6	Twenty-one routed modes in the routing table; alignment is the required mode	skills/helix/SKILL.md#routing-rules , docs/helix/01-frame/prd.md#goals

Sources (continued)

Id	Claim or figure	Governing artifact and section
S7	Four-step contract: write the brief, check alignment, plan the work, run it	docs/helix/01-frame/prd.md#core-workflow-contract
S8	OAuth example: three specs and two designs affected, one decision in conflict, four stories missing	docs/helix/00-discover/product-vision.md#user-experience
S9	Concerns propagate practices downstream; the project's active concerns and slots; 50 concerns in the library	workflows/README.md#cross-cutting-context, docs/helix/01-frame/concerns.md, workflows/concerns
S10	Seven principles, from spec is the contract to deliverable over machinery, with tension resolution	workflows/principles.md#principles, docs/helix/01-frame/principles.md
S11	Artifact honesty: phantom claims block; verification evidence comes from the running system	docs/helix/01-frame/features/FEAT-016-artifact-honesty.md, workflows/concerns/verification/practices.md
S12	Tests cite the criterion they cover; quality floors never slide back	docs/helix/02-design/adr/ADR-009-acceptance-criteria-ownership.md, workflows/ratchets.md

Sources (continued)

Id	Claim or figure	Governing artifact and section
S13	Non-goals: no CLI, tracker, queue, or engine; no imposed technology; judgment stays with people	docs/helix/01-frame/prd.md#non-goals
S14	Six install surfaces: DDx, Claude Code, Codex, Copilot, Databricks Genie, Grok Build	docs/helix/01-frame/features/FEAT-013-runtime-install-coverage.md, docs/install/README.md
S15	Autonomy is a three-position spectrum with a hard-stop floor at every level	docs/helix/02-design/adr/ADR-003-autonomy-spectrum.md
S16	Roadmap, iteration plan, status report; a client cut may subset but never contradict	docs/helix/01-frame/features/FEAT-017-iteration-documentation.md
S17	Deliverables projected from governed artifacts, every number sourced, gated and inspected	workflows/modes/present.md, workflows/activities/06-iterate/artifacts/deliverable/meta.yml
S18	Targets: < 3 findings per run; 3 runtimes in 12 months; 50+ repos in 18 months	docs/helix/00-discover/product-vision.md#success-definition, docs/helix/01-frame/prd.md#success-metrics